

Using relevance feedback and text similarity to reduce review effort in eDiscovery

Supervisors:

Dr. P.W.H vd Putten, Prof. Dr. Ir. J.C. Scholtes
L. Jean-Louis PhD & Z. Gerolemou MSc

Timo Kats, Universiteit Leiden.



**Universiteit
Leiden**
The Netherlands

Discover the world at Leiden University

Using text similarity and relevance feedback to reduce review effort

eDiscovery context 1/2

1 | 21

- Prior to lawsuits and investigations there's a *discovery* process.
- Produce *evidence* to the other party.
- *eDiscovery* refers to this process, but with the inclusion of electronically stored information (e.g. emails, PDFs).
- A set of documents is first collected and thereafter *manually reviewed*.

- The *manual reviewing* of the collected evidence is the biggest cost driver in eDiscovery.
- High *review effort* (i.e. the amount of documents that have to be reviewed), is the main driver of these costs.

The Loneliness of the Long-Distance Document Reviewer: E-discovery and Cognitive Ergonomics [ADB09]

"To fulfill a similar transaction, the firm [Crowell & Moring] employed 125 contract lawyers for three months. They reviewed 30 million pages and produced 12 million relevant pages."

Problems in document review

3 | 21

- There are methods that reduce review effort.
- An **active learning based approach** with a human in the loop (also referred to as “Technology Assisted Review”)
- Although effective at reducing review effort, a possibility of **false negatives** without the awareness of the user.

Tarbits: The sticky consequences of poorly implementing technology assisted review [Dow20]

“90% recall of relevant documents misses that the 10% of false negatives may be not just relevant, but crucial, even to the point of being more worthwhile than the other 90% altogether”

Research goal and questions

4 | 21

Research goal

- Use relevance feedback to reduce review effort.
 - **Relevance feedback** refers to the iterative process of improving the **relevance ranking** based on user feedback.
 - The **relevance ranking** will be based on the **textual similarity** between a queried document and the remaining documents.

Research questions

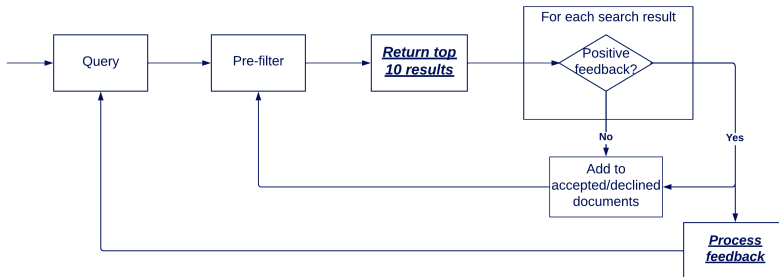
1. Can text similarity be used for relevance feedback?
2. To what extent can relevance feedback help to reduce review effort?

- RCV 1 v2 dataset (annotated Reuters news articles)
- Test set (topics fairly unrelated to each other)
 - Strategy/Plans
 - Regulation/Policy
 - War/Civil War
 - Sports
 - Elections
- “Ambiguous” set (same parent topics)
 - EC Corporate policy
 - EC Internal market
 - Forex markets
 - Energy markets
- Random sample of 300 articles per topic.

Using text similarity and relevance feedback to reduce review effort

Method

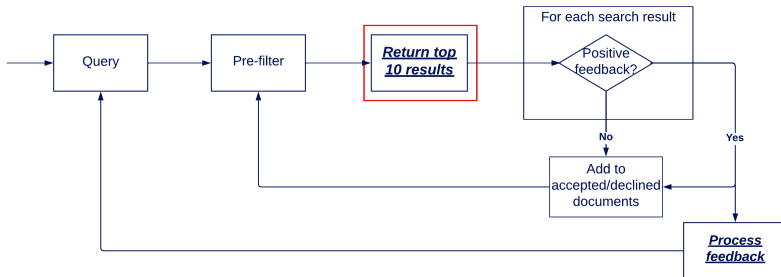
6 | 21



Using text similarity and relevance feedback to reduce review effort

Method

6 | 21



Using text similarity and relevance feedback to reduce review effort

Text similarity - Components

7 | 21

Pre-processing

- Removing numbers, special characters and text to lowercase.

Creating embeddings

- BERT based embeddings

Ranking

- Cosine distance between the embeddings

Implementation

- Solr's Dense Vector Search (DVS)

Using text similarity and relevance feedback to reduce review effort

Text similarity - Embeddings

8 | 21

- *BERT* creates dense word embeddings that capture a form of *semantic meaning* through context.
- *SBERT* creates paragraph level embeddings through *mean pooling* these word embeddings (no document level).

Timo is a student > [-1.8472e-01, -3.1975e-01, 2.0524e-01, ...]

SBERT [RG19]

“This reduces the effort for finding the most similar pair in a collection of 10,000 sentences from **65 hours** with BERT / RoBERTa to about **5 seconds** with SBERT...”

Text similarity - Ranking

9 | 21

- Text similarity can be computed through the *cosine distance* between the embeddings.
- On the paragraph level, the same document can be returned multiple times.

Paragraph based document rankings

Return the following 6 paragraphs: $\{d_1, d_2, d_2, d_3, d_3, d_3\}$
(where d_i refers to a paragraph from document i).

The first based ranking would be: $\{1, 2, 3\}$ whereas the count based ranking would be: $\{3, 2, 1\}$

Using text similarity and relevance feedback to reduce review effort

Results - DVS configurations

10 | 21

Recall-Precision graph shows that DVS using the **first based ranking** outperforms DVS based on the **count based ranking**.





Figure: *DVS configurations*

Using text similarity and relevance feedback to reduce review effort


Results - Baselines

11 | 21

Recall-Precision graph show that this **DVS** configuration outperforms the **Quorum** (occurrence only) and **TF-IDF** based approaches.



(a) Test set



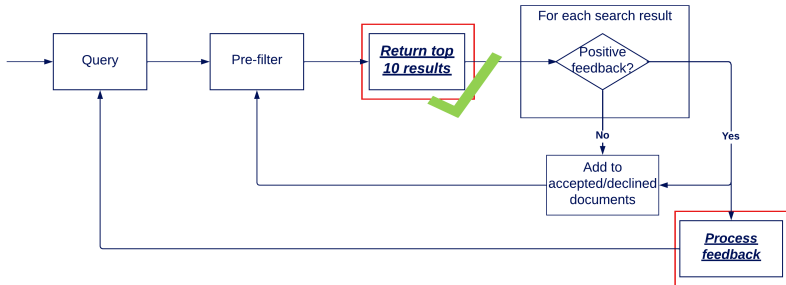
(b) Ambiguous set

Figure: Different configurations for DVS

Using text similarity and relevance feedback to reduce review effort

Method

12 | 21



Relevance feedback - Components

13 | 21

Pseudo-relevance feedback

- *Improve* the queried embedding based on the *selected* search results (which re-ranks the results for the next ranking).
- We *assume* that articles that contain the queried topic get *positive feedback* from the “user”.

Feedback strategies

- Commonly used strategies/baselines:
 - No feedback (show next 10 results)
 - Keyword expansion
 - Rocchio
- Average/Sum the SBERT embeddings

Relevance feedback – Baselines

14 | 21

Keyword expansion, each iteration:

1. Get words from *selected* texts
2. Sort words by their *IDF* value (uniqueness)
3. Append the top 10 words to *keyword filter*
4. Pre-filter results with: keyword1 OR keyword2 OR ...

Rocchio, each iteration:


1. Get average embedding of selected texts (α).
2. Average this (weighted) with queried embedding (β).
3. For this weighted average, we use $\alpha = 0.5, \beta = 0.5$

Using text similarity and relevance feedback to reduce review effort

Relevance feedback - Strategies

15 | 21


Based on averaging/summing the SBERT embeddings. Variations based on (1) cumulative feedback and (2) feedback amplification.




Using text similarity and relevance feedback to reduce review effort

Relevance feedback – Review effort 16 | 21

- **Cumulatively** summing the embeddings reduces review effort the most (measured in iterations needed to achieve 80% recall).
- Adding sibling paragraphs reduces the standard deviation.



(a) Test set



(b) Ambiguous set

Figure: Results of the relevance feedback strategies

Using text similarity and relevance feedback to reduce review effort

Relevance feedback - Latency

17 | 21

As for efficiency, summing/averaging the embeddings add little to no latency. Amplifying feedback does.




Figure: Average iteration times for different feedback strategies

Findings

18 | 21

- Compared to our minimal baseline (of no feedback), review effort is reduced between 17.85% (on the test set) and 59.04% (on the ambiguous set).
- Compared to an SVM based approach, this method is very fast.
- This method doesn't produce false negatives without the awareness of the user.

Limitations and future work

19 | 21

Limitations:

- Data size and data homogeneity.

Future work suggestions:

- Heterogeneous data (emails, contracts, direct messages, etc.)
- Chunking (compute text similarities in parallel)
- Generative AI
 - Help formulating queries
 - Summarizing results
 - A lot of developments that can be *complementary* to our method.

Using text similarity and relevance feedback to reduce review effort

In conclusion

20 | 21

We found that...

- Our method reduced review effort between 17.85% and 59.04%.
- Very little latency.

However...

- Data homogeneity and size unrepresentative of eDiscovery

Regardless...

- Documents are only re-ranked, so no false negatives without the awareness of the user.
- Method applicable in real-world eDiscovery scenarios.

Using text similarity and relevance feedback to reduce review effort

Thank you

21 | 21

Any questions?

```
#include <iostream>
#include <string>

int main()
{
    std::string supervisors[4] = {"Jan", "Peter", "Zoe", "Ludovic"};

    for(int i = 0; i <= supervisors->length(); i++)
    {
        std::cout << "thank you " << supervisors[i] << std::endl;
    }
    return 0;
}
```

References

21 | 21

- [ADB09] Simon Attfield, Stephen De Gabrielle, and Ann Blandford. “The loneliness of the long-distance document reviewer: e-Discovery and cognitive ergonomics”. In: *DESI III workshop at ICAIL, Barcelona*. Citeseer. 2009.
- [RG19] Nils Reimers and Iryna Gurevych. “Sentence-bert: Sentence embeddings using siamese bert-networks”. In: *arXiv preprint arXiv:1908.10084* (2019).
- [Dow20] David Dowling. “Tarpits: The Sticky Consequences of Poorly Implementing Technology-Assisted Review”. In: *Berkeley Tech. LJ* 35 (2020), p. 171.

Using text similarity and relevance feedback to reduce review effort

Additional information


21 | 21

Using text similarity and relevance feedback to reduce review effort


TF-IDF/Quorum - paragraph level

21 | 21

For both methods the **document level** outperforms both paragraph based rankings.



(a) *TF-IDF*



(b) *Quorum search*

Figure: *Quorum and TF-IDF on the document/paragraph level*

Implementing this method

21 | 21

- DVS can be implemented out-of-the-box when using search engines like *Solr*¹.
- SBERT embeddings can be created using the *sentence-transformers*² Python library.
- The relevance feedback is implemented using *Numpy*³.

¹<https://solr.apache.org/>

²<https://www.sbert.net/>

³<https://numpy.org/>

Using text similarity and relevance feedback to reduce review effort

Error analysis

21 | 21

Manually sampled and reviewed 50FPs and 50FNs per method and segmented into buckets. For the identified method (DVS):

False positives

- 10 due to “non topic similarities”
- 33 due to “partial topic overlap”
- 7 due to “bad paragraph segmentation”

False negatives

- 25 due to “dissimilar content”
- 9 due to “bad paragraph segmentation”
- 16 paragraphs are “too specific to match”

Using text similarity and relevance feedback to reduce review effort

DVS optimization

21 | 21

DVS doesn't use chunking. Instead, it uses HNSW, which finds closest neighbors through (greedily) traversing through an hierarchical NSW network graph.




Figure: *Illustration of the HNSW algorithm*